

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_iris
from sklearn.metrics import classification_report

%matplotlib inline
```

本文主要参考了以下材料：

1. cs229: 9.3 softmax regression
2. <http://ufldl.stanford.edu/wiki/index.php/Softmax%E5%9B%9E%E5%BD%92>
(<http://ufldl.stanford.edu/wiki/index.php/Softmax%E5%9B%9E%E5%BD%92>).
3. <https://blog.csdn.net/u012328159/article/details/72155874>
(<https://blog.csdn.net/u012328159/article/details/72155874>).

Softmax Regression

1. 原理推导

Consider a classification problem in which the response variable y can take on any one of k values, so $y \in \{1, 2, \dots, k\}$. The response variable is still discrete, but can now take on more than two values. We will thus model it as distributed according to a multinomial distribution.

Lets derive a GLM for modelling this type of multinomial data. To do so, we will begin by expressing the multinomial as an exponential family distribution.

To parameterize a multinomial over k possible outcomes, one could use k parameters ϕ_1, \dots, ϕ_k specifying the probability of each of the outcomes. However, these parameters would be redundant, or more formally, they

would not be independent (since knowing any $k - 1$ of the ϕ_i 's uniquely determines the last one, as they must satisfy $\sum_{i=1}^k \phi_i = 1$). So, we will instead parameterize the multinomial with only $k - 1$ parameters, $\phi_1, \dots, \phi_{k-1}$, where $\phi_i = p(y = i; \phi)$, and $p(y = k; \phi) = 1 - \sum_{i=1}^{k-1} \phi_i$. For notational convenience, we will also let $\phi_k = 1 - \sum_{i=1}^{k-1} \phi_i$, but we should keep in mind that this is not a parameter, and that it is fully specified by $\phi_1, \dots, \phi_{k-1}$.

To express the multinomial as an exponential family distribution, we will define $T(y) \in \mathbb{R}^{k-1}$ as follows:

$$T(1) = \begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, T(2) = \begin{bmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, T(3) = \begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}, \dots, T(k-1) = \begin{bmatrix} 0 \\ 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}, T(k) = \begin{bmatrix} 0 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix},$$

Unlike our previous examples, here we do not have $T(y) = y$; also, $T(y)$ is now a $k - 1$ dimensional vector, rather than a real number. We will write $(T(y))_i$ to denote the i -th element of the vector $T(y)$.

We introduce one more very useful piece of notation. An indicator function $1\{\cdot\}$ takes on a value of 1 if its argument is true, and 0 otherwise. So, we can write the relationship between $T(y)$ and y as $(T(y))_i = 1\{y = i\}$ (当且仅当 $y = i$ 时, 向量 $T(y)$ 的第 i 个位置元素为 1). Further, we have that $E[(T(y))_i] = P(y = i) = \phi_i$.

We are now ready to show that the multinomial is a member of the exponential family. We have:

$$\begin{aligned}
p(y; \phi) &= \phi_1^{1\{y=1\}} \phi_2^{1\{y=2\}} \dots \phi_k^{1\{y=k\}} \\
&= \phi_1^{1\{y=1\}} \phi_2^{1\{y=1\}} \dots \phi_k^{1 - \sum_{i=1}^{k-1} 1\{y=i\}} \\
&= \phi_1^{(T(y))_1} \phi_2^{(T(y))_2} \dots \phi_k^{1 - \sum_{i=1}^{k-1} (T(y))_i} \\
&= \exp((T(y))_1 \log(\phi_1) + (T(y))_2 \log(\phi_2) + \dots + (1 - \sum_{i=1}^{k-1} (T(y))_i) \log(\phi_k)) \\
&= \exp((T(y))_1 \log(\frac{\phi_1}{\phi_k}) + (T(y))_2 \log(\frac{\phi_2}{\phi_k}) + \dots + (T(y))_{k-1} \log(\frac{\phi_{k-1}}{\phi_k}) + \log(\phi_k)) \\
&= b(y) \exp(\eta^T T(y) - a(\eta))
\end{aligned}$$

where

$$\eta = \begin{bmatrix} \log(\frac{\phi_1}{\phi_k}) \\ \log(\frac{\phi_2}{\phi_k}) \\ \vdots \\ \log(\frac{\phi_{k-1}}{\phi_k}) \end{bmatrix}$$

$$a(\eta) = -\log(\phi_k)$$

$$b(y) = 1$$

This completes our formulation of the multinomial as an exponential family distribution.

The link function is given (for $i = 1, \dots, k$) by

$$\eta_i = \log \frac{\phi_i}{\phi_k}$$

For convenience, we have also defined $\eta_k = \log(\frac{\phi_k}{\phi_k}) = 0$. To invert the link function and derive the response function, we therefore have that

$$e^{\eta_i} = \frac{\phi_i}{\phi_k}$$

$$\phi_k e^{\eta_i} = \phi_i \quad - (1)$$

$$\phi_k \sum_{i=1}^k e^{\eta_i} = \sum_{i=1}^k \phi_i = 1$$

This implies that $\phi_k = 1/\sum_{i=1}^k e^{\eta_i}$, which can be substituted back into Equation (1) to give the response function

$$\phi_i = \frac{e^{\eta_i}}{\sum_{l=1}^k e^{\eta_l}}$$

This function mapping from the η 's to the ϕ 's is called the **softmax function**.

To complete our model, we use Assumption 3, given earlier, that the η 's are linearly related to the x 's. So, have $\eta_i = \theta_i^T x$ (for $i = 1, \dots, k-1$), where $\theta_1, \dots, \theta_{k-1} \in \mathbb{R}^{n+1}$ are the parameters of our model. For notational convenience, we can also define $\theta_k = 0$, so that $\eta_k = \theta_k^T x = 0$, as given previously. Hence, our model assumes that the conditional distribution of y given x is given by

$$\begin{aligned}
 p(y = i|x; \theta) &= \phi_i \\
 &= \frac{e^{\eta_i}}{\sum_{l=1}^k e^{\eta_l}} \\
 &= \frac{e^{\theta_i^T x}}{\sum_{l=1}^k e^{\theta_l^T x}} \quad - (2)
 \end{aligned}$$

This model, which applies to classification problems where $y \in \{1, \dots, k\}$, is called **softmax regression**. It is a generalization of logistic regression.

Our hypothesis will output

$$\begin{aligned}
h_{\theta}(x) &= E[T(y) | x; \theta] \\
&= E \left[\begin{array}{c|c} \begin{matrix} 1\{y = 1\} \\ 1\{y = 2\} \\ \dots \\ 1\{y = k-1\} \end{matrix} & x; \theta \end{array} \right] \\
&= \begin{bmatrix} \phi_1 \\ \phi_2 \\ \dots \\ \phi_{k-1} \end{bmatrix} \\
&= \begin{bmatrix} \frac{\exp(\theta_1^T x)}{\sum_{l=1}^k \exp(\theta_l^T x)} \\ \frac{\exp(\theta_2^T x)}{\sum_{l=1}^k \exp(\theta_l^T x)} \\ \vdots \\ \frac{\exp(\theta_{k-1}^T x)}{\sum_{l=1}^k \exp(\theta_l^T x)} \end{bmatrix}
\end{aligned}$$

In other words, our hypothesis will output the estimated probability that $p(y = i|x; \theta)$, for every value of $i = 1, \dots, k$. (Even though $h_{\theta}(x)$ as defined above is only $k-1$ dimensional, clearly $p(y = k|x; \theta)$ can be obtained as $1 - \sum_{i=1}^{k-1} \phi_i$.)

Lastly, let's discuss parameter fitting. Similar to our original derivation of ordinary least squares and logistic regression, if we have a training set of m examples $\{(x^{(i)}, y^{(i)}); i = 1, \dots, m\}$ and would like to learn the parameter θ_i of this model, we would begin by writing down the log-likelihood:

$$\begin{aligned}
l(\theta) &= \sum_{i=1}^m \log p(y^{(i)} | x^{(i)}; \theta) \\
&= \sum_{i=1}^m \log \prod_{j=1}^k \left(\frac{e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}} \right)^{1_{\{y^{(i)}=j\}}} \\
&= \sum_{i=1}^m \sum_{j=1}^k 1_{\{y^{(i)}=j\}} \log \frac{e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}}
\end{aligned}$$

To obtain the second line above, we used the definition for $p(y|x; \theta)$ given in Equation (2). We can now obtain the maximum likelihood estimate of the parameters by maximizing $l(\theta)$ in terms of θ , using a method such as gradient ascent or Newton's method.

2. 梯度下降法

2.1 Cost Function

首先，定义

$$x^{(i)} = \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ \vdots \\ x_{n+1}^{(i)} \end{bmatrix}_{(n+1) \times 1} \quad X = \begin{bmatrix} -(x^{(1)})^T - \\ -(x^{(2)})^T - \\ \dots \\ -(x^{(m)})^T - \end{bmatrix}_{m \times (n+1)} \quad \theta = \begin{bmatrix} -\theta_1^T - \\ -\theta_2^T - \\ \dots \\ -\theta_k^T - \end{bmatrix}_{k \times (n+1)}$$

注意到， $x^{(i)} \in \mathbb{R}^{n+1}$ ，其中定义 $x_0^{(i)} = 0$ 。然后， $\theta_1, \theta_2, \dots, \theta_k \in \mathbb{R}^{n+1}$ ，在这里，我没有定义 $\theta_k = 0$ ，因

此，存在过度参数化的问题。所以，在接下来的小节内会对这个问题进行解决。最后， $y \in \{1, 2, \dots, k\}$ 。

因此，我们定义Cost Function为

$$J(\theta) = -\frac{1}{m} \left(\sum_{i=1}^m \sum_{j=1}^k 1\{y^{(i)} = j\} \log \frac{e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}} \right)$$

值得注意的是，上述公式是逻辑回归的cost function的推广，逻辑回归的cost function可以改为：

$$\begin{aligned} J(\theta) &= -\frac{1}{m} \left(\sum_{i=1}^m (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) + y^{(i)} \log h_{\theta}(x^{(i)}) \right) \\ &= -\frac{1}{m} \left(\sum_{i=1}^m \sum_{j=0}^1 1\{y^{(i)} = j\} \log p(y^{(i)} = j | x^{(i)}; \theta) \right) \end{aligned}$$

2.2 Softmax Regression 模型参数化的特点

由于上一节，我们没有定义 $\theta_k = 0$ ，这使得softmax regression有一个“冗余”的参数集。虽然，定义 $\theta_k = 0$ 可以避免这个问题，但是这会使得在算法实现中没有那么简单清楚，而且，这个问题是可以得到解决的。接下来，我们对这个问题进行具体说明。

假设我们从参数向量 θ_j 中减去了向量 ψ ，这时，每一个 θ_j 都变成了 $\theta_j - \psi$ ($j = 1, \dots, k$)。此时假设函数变成了以下的式子：

$$\begin{aligned}
p(y^{(i)} = j | x^{(i)}; \theta) &= \frac{e^{(\theta_j - \psi)^T x^{(i)}}}{\sum_{l=1}^k e^{(\theta_l - \psi)^T x^{(i)}}} \\
&= \frac{e^{\theta_j^T x^{(i)}} e^{-\psi^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}} e^{-\psi^T x^{(i)}}} \\
&= \frac{e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}}
\end{aligned}$$

换句话说，从 θ_j 中减去 ψ 完全不影响假设函数的预测结果。这表明前面的 softmax regression 模型中存在冗余的参数。更正式一点说，softmax regression 模型被过度参数化了。对于任意一个用于拟合数据的假设函数，可以求出多组参数值，这些参数得到的是完全相同的假设函数 h_θ 。

进一步而言，如果参数 $(\theta_1, \theta_2, \dots, \theta_k)$ 是 cost function $J(\theta)$ 的极小值点，那么 $(\theta_1 - \psi, \theta_2 - \psi, \dots, \theta_k - \psi)$ 同样也是它的极小值点，其中 ψ 可以为任意向量。因此，使得 $J(\theta)$ 最小化的解不是唯一的。（有趣的是，由于 $J(\theta)$ 仍然是一个凸函数，因此梯度下降不会遇到局部最优解的问题，但是 *Hessian* 矩阵是奇异的/不可逆的，这会直接导致采用牛顿法优化就遇到数值计算的问题）

注意，当 $\phi = \theta_k$ 时，我们总是可以将 θ_k 替换为 $\theta_k - \psi = \vec{0}$ （即替换为全零向量），并且这种替换不会影响假设函数。因此，我们可以去掉参数向量 θ_k （或者其他 θ_j 中的任意一个）而不影响假设函数的表达能力。实际上，与其优化全部的 $k \times (n + 1)$ 个参数 $(\theta_1, \theta_2, \dots, \theta_k)$ （其中， $\theta_j \in \mathbb{R}^{n+1}$ ），我们可以令 $\theta_k = \vec{0}$ ，只优化剩余的 $(k - 1) \times (n + 1)$ 个参数，这样算法依然能够正常工作。

在实际应用中，为了让算法实现更加简单清楚，往往保留所有参数 $(\theta_1, \theta_2, \dots, \theta_k)$ ，而不任意地将某一参数设置为 $\vec{0}$ 。但此时，我们需要对 cost function 做一个改动：加入权重衰减项。权重衰减项可以解决 softmax regression 参数冗余所带来的数值问题。

2.3 权重衰减

我们通过添加一个权重衰减项 $\frac{\lambda}{2} \sum_{i=1}^k \sum_{j=0}^n \theta_{ij}^2$ 来修改 cost function，这个衰减项会惩罚过大的参数值，现在我们的 cost function 变为：

$$J(\theta) = -\frac{1}{m} \left(\sum_{i=1}^m \sum_{j=1}^k 1\{y^{(i)} = j\} \log \frac{e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}} \right) + \frac{\lambda}{2} \sum_{i=1}^k \sum_{j=0}^n \theta_{ij}^2$$

有了这个权重衰减项之后($\lambda > 0$)，Cost function 就变成了严格的凸函数，这样就可以保证得到唯一的解了。此时的 Hessian 矩阵变成可逆矩阵，并且因为 $J(\theta)$ 是凸函数，梯度下降法和L-BFGS等算法可以保证收敛到全局最优解。

为了使用优化算法，我们需要求得这个新函数的 $J(\theta)$ 的导数，如下：

$$\nabla_{\theta_j} J(\theta) = -\frac{1}{m} \left(x^{(i)} (1\{y^{(i)} = j\} - p(y^{(i)} | x^{(i)}; \theta)) \right) + \lambda \theta_j$$

注意，这里的 $\nabla_{\theta_j} J(\theta) \in \mathbb{R}^{n+1}$ 。通过最小化 $J(\theta)$ ，我们就能实现一个可用的 softmax regression 模型。

2.4 推导 $\frac{\partial J(\theta)}{\partial \theta_j}$

方法一：

$$\begin{aligned}
\frac{\partial J(\theta)}{\partial \theta_j} &= -\frac{1}{m} \frac{\partial}{\partial \theta_j} \left[\sum_{i=1}^m \sum_{j=1}^k 1\{y^{(i)} = j\} \log \frac{e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}} \right] + \lambda \theta_j \\
&= -\frac{1}{m} \frac{\partial}{\partial \theta_j} \left[\sum_{i=1}^m \sum_{j=1}^k 1\{y^{(i)} = j\} (\log e^{\theta_j^T x^{(i)}} - \log \sum_{l=1}^k e^{\theta_l^T x^{(i)}}) \right] + \lambda \theta_j \\
&= -\frac{1}{m} \frac{\partial}{\partial \theta_j} \left[\sum_{i=1}^m \sum_{j=1}^k 1\{y^{(i)} = j\} (\theta_j^T x^{(i)} - \log \sum_{l=1}^k e^{\theta_l^T x^{(i)}}) \right] + \lambda \theta_j \\
&= -\frac{1}{m} \frac{\partial}{\partial \theta_j} \left[\sum_{i=1}^m 1\{y^{(i)} = j\} \left(\sum_{j=1}^k \theta_j^T x^{(i)} - \sum_{j=1}^k \log \sum_{l=1}^k e^{\theta_l^T x^{(i)}} \right) \right] + \lambda \theta_j \\
&= -\frac{1}{m} \left[\sum_{i=1}^m 1\{y^{(i)} = j\} \left(x^{(i)} - \sum_{j=1}^k \frac{x^{(i)} e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}} \right) \right] + \lambda \theta_j \\
&= -\frac{1}{m} \left[\sum_{i=1}^m x^{(i)} (1\{y^{(i)} = j\} - \sum_{j=1}^k 1\{y^{(i)} = j\} \frac{e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}}) \right] + \lambda \theta_j \\
&= -\frac{1}{m} \left[\sum_{i=1}^m x^{(i)} \left(1\{y^{(i)} = j\} - \frac{e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}} \right) \right] + \lambda \theta_j \\
&= -\frac{1}{m} \left[\sum_{i=1}^m x^{(i)} (1\{y^{(i)} = j\} - p(y^{(i)} | x^{(i)}; \theta)) \right] + \lambda \theta_j
\end{aligned}$$

方法二：

$$\begin{aligned}
\frac{\partial J(\theta)}{\partial \theta_j} &= -\frac{1}{m} \left[\sum_{i=1}^m \frac{\partial}{\partial \theta_j} (1\{y^{(i)} = j\} \log \frac{e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}} + \sum_{c \neq j}^k 1\{y^{(i)} = c\} \log \frac{e^{\theta_c^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}}) \right] + \lambda \theta_j \\
&= -\frac{1}{m} \left[\sum_{i=1}^m (1\{y^{(i)} = j\} (x^{(i)} - \frac{x^{(i)} e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}}) + \sum_{c \neq j}^k 1\{y^{(i)} = c\} (-\frac{x^{(i)} e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}})) \right] + \lambda \theta_j \\
&= -\frac{1}{m} \left[\sum_{i=1}^m (x^{(i)} 1\{y^{(i)} = j\} (1 - \frac{e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}}) + \sum_{c \neq j}^k 1\{y^{(i)} = c\} (-\frac{x^{(i)} e^{\theta_j^T x^{(i)}}}{\sum_{l=1}^k e^{\theta_l^T x^{(i)}}})) \right] + \lambda \theta_j \\
&= -\frac{1}{m} \left[\sum_{i=1}^m x^{(i)} (1\{y^{(i)} = j\} - 1\{y^{(i)} = j\} p(y^{(i)} | x^{(i)}; \theta)) + \sum_{c \neq j}^k 1\{y^{(i)} = c\} (-p(y^{(i)} | x^{(i)}; \theta)) \right] + \lambda \theta_j \\
&= -\frac{1}{m} \left[\sum_{i=1}^m x^{(i)} (1\{y^{(i)} = j\} - \sum_{j=1}^k 1\{y^{(i)} = j\} p(y^{(i)} | x^{(i)}; \theta)) \right] + \lambda \theta_j \\
&= -\frac{1}{m} \left[\sum_{i=1}^m x^{(i)} (1\{y^{(i)} = j\} - p(y^{(i)} | x^{(i)}; \theta)) \right] + \lambda \theta_j
\end{aligned}$$

2.5 矩阵化

因为 $y \in \{1, 2, \dots, k\}$, 所以, 对 y 进行独热编码。即 $y^{(i)} \in \mathbb{R}^k$, 其中, 若 $y^{(i)}$ 属于类别 i , 则 $y^{(i)}$ 第 i 个位置上的元素为1, 其余位置元素为0。因此, 定义矩阵 G 为

$$G = \begin{bmatrix} -(y^{(1)})^T - \\ -(y^{(2)})^T - \\ \vdots \\ -(y^{(m)})^T - \end{bmatrix}_{m \times k}$$

定义概率矩阵 P 为

$$P_{m \times k} = \text{norm}(\exp(X_{m \times (n+1)} \cdot \theta_{(n+1) \times k}^T))$$

其中，norm 表示归一化项，因此，概率矩阵 P 的具体计算方式为：首先，计算 $\exp(X\theta^T)$ 得到 $m \times k$ 的矩阵。其次，使用 `np.sum()` 对该矩阵按行进行求和，得到 $m \times 1$ 的矩阵。最后，利用Python的广播（broadcast）机制，将该矩阵与 $\exp(X\theta^T)$ 对应位置元素进行相乘（element-wise multiplication）得到概率矩阵 P 。

于是

$$\frac{\partial J(\theta)}{\partial \theta} = -\frac{1}{m}(G - P)^T \cdot X + \lambda \theta$$

所以，cost function为

$$J(\theta) = -np.mean(G \circ P) + \lambda np.sum(\theta)$$

其中， \circ 表示对应位置元素相乘，即 element-wise multiplication。

3. 实现 softmax regression

3.1 读取数据

```
In [2]: iris = load_iris()
features = pd.DataFrame(data=iris.data, columns=iris.feature_names)
label = pd.DataFrame(data=iris.target, columns=['traget'])
data = pd.concat([features, label], axis=1)
data.head()
```

Out[2]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	traget
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0


```
In [3]: def loadData(df):
ones = pd.DataFrame({'ones': np.ones(len(df))})
df = pd.concat([ones, df], axis=1)
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
return X, y
```

```
In [4]: X, y = loadData(data)
X.shape, y.shape
```

Out[4]: ((150, 5), (150,))

3.2 数据可视化

```
In [5]: plt.figure(figsize=(12, 8))
plt.scatter(X[:, 3][y==0], X[:, 4][y==0], marker='x', label=iris.target_names[0])
plt.scatter(X[:, 3][y==1], X[:, 4][y==1], marker='^', label=iris.target_names[1])
plt.scatter(X[:, 3][y==2], X[:, 4][y==2], marker='+', label=iris.target_names[2])
plt.legend(loc='upper left')
plt.xlabel(iris.feature_names[2])
plt.ylabel(iris.feature_names[3])
plt.title('Iris Data Set')
plt.show()
```



3.3 Softmax Regression

对 y 进行独热编码，得到矩阵 G 。

```
In [6]: def oneHotY(y):  
        # m为样本数  
        m = y.shape[0]  
        # k为类别数  
        k = len(np.unique(y))  
  
        oneHotY = np.zeros((m, k))  
        for i in range(k):  
            oneHotY[:, i] = (y==i)  
  
        return oneHotY
```

```
In [7]: G = oneHotY(y)  
        G.shape
```

```
Out[7]: (150, 3)
```

```
In [8]: def initializeWithZeros(X, y):  
        k = len(np.unique(y))  
  
        return np.zeros((k, X.shape[1]))
```

```
In [9]: def probabilityMatrix(X, theta):  
        expScore = np.exp(X @ theta.T)  
        sumScore = np.sum(expScore, axis=1).reshape(-1, 1)  
  
        return np.multiply(expScore, 1 / sumScore)
```

```
In [10]: def computeCost(X, G, theta, l):  
        P = probabilityMatrix(X, theta)  
        return -np.mean(np.multiply(G, np.log(P))) + l * theta.sum()
```



```
In [11]: def computeGradient(X, G, theta, l):  
    m = X.shape[0]  
    P = probabilityMatrix(X, theta)  
    grad = -((G-P).T @ X) / m + l * theta  
  
    return grad
```

```
In [12]: def batchGradientDescent(X, G, theta, alpha, iters, l, printFlag=True):  
    costs = np.zeros(iters)  
  
    for i in range(iters):  
        theta = theta - alpha * computeGradient(X, G, theta, l)  
        costs[i] = computeCost(X, G, theta, l)  
  
        if printFlag and i % 1000 == 0:  
            print(costs[i])  
  
    return theta, costs
```

```
In [13]: def predict(X, theta):  
    P = probabilityMatrix(X, theta)  
  
    return np.argmax(P, axis=1).reshape(-1, 1)
```

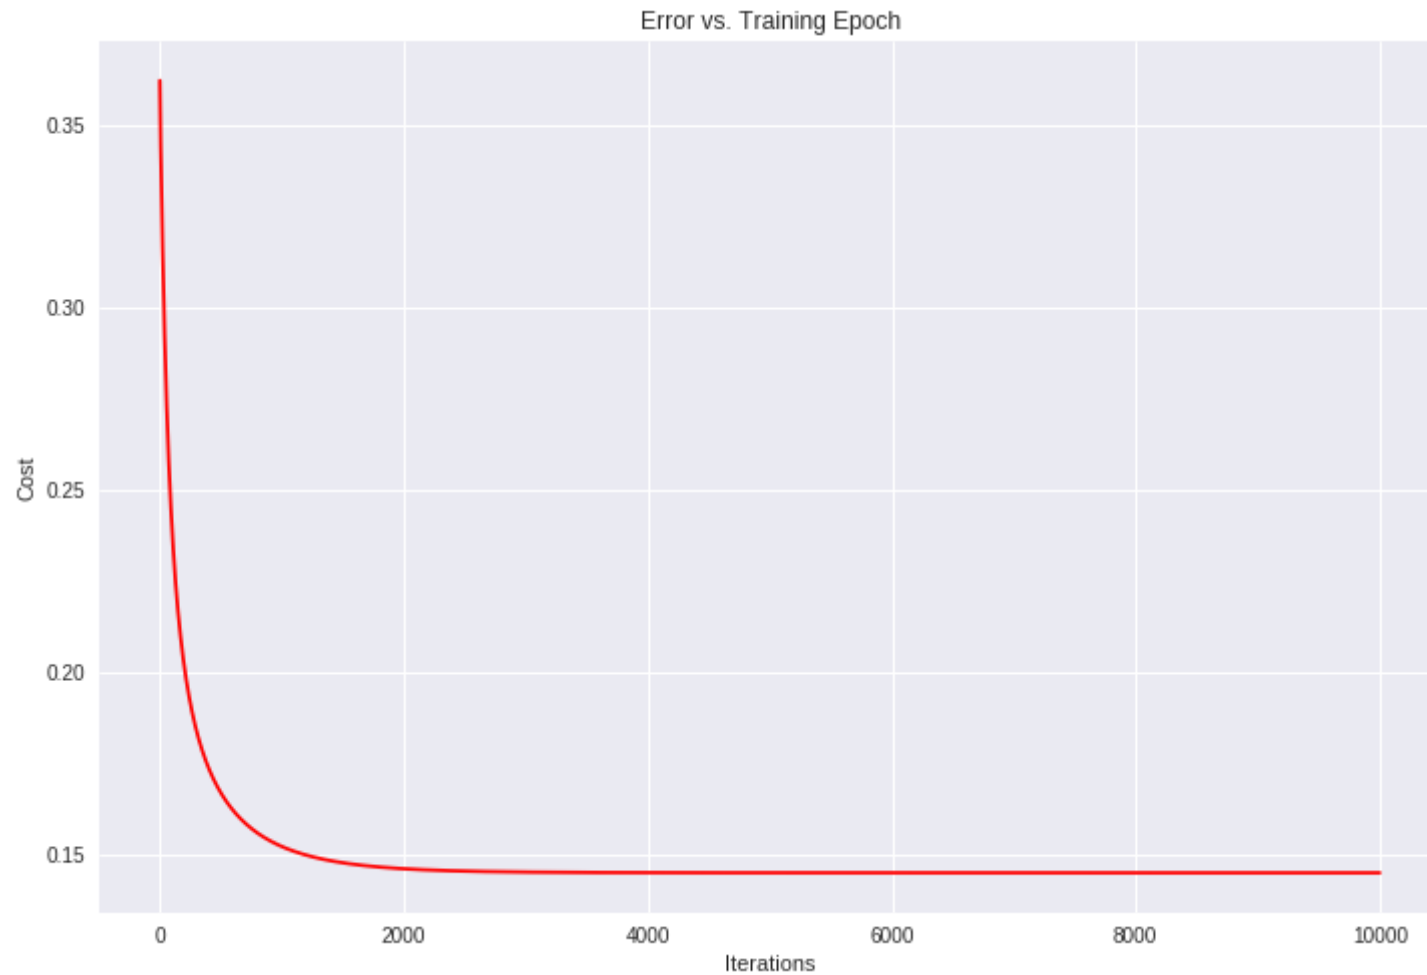
```
In [14]: X, y = loadData(data)
G = oneHotY(y)
theta = initializeWithZeros(X, y)

iters = 10000
alpha = 0.01
l = 0.1

theta, costs = batchGradientDescent(X, G, theta, alpha, iters, l)
```

```
0.36204899345585234
0.15207044998459318
0.14597815071859724
0.14507033494591312
0.14491252983436856
0.14488189983190422
0.14487517819681003
0.14487351365380918
0.14487305861140318
0.144872925162167
```

```
In [15]: plt.figure(figsize=(12, 8))  
plt.plot(np.arange(iters), costs, color='red')  
plt.xlabel('Iterations')  
plt.ylabel('Cost')  
plt.title('Error vs. Training Epoch')  
plt.show()
```



```
In [16]: y_pred = predict(X, theta)
print(classification_report(y, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	50
1	1.00	0.88	0.94	50
2	0.89	1.00	0.94	50
accuracy			0.96	150
macro avg	0.96	0.96	0.96	150
weighted avg	0.96	0.96	0.96	150

4. Softmax Regression 与 Logistic Regression 的关系

当类别数 $k = 2$ 时，softmax regression 退化为 logistic regression。这表明 softmax regression 是 logistic regression 的一般形式。具体地说，当 $k = 2$ 时，softmax regression 的假设函数为：

$$h_{\theta}(x) = \frac{1}{e^{\theta_1^T x} + e^{\theta_2^T x}} \begin{bmatrix} e^{\theta_1^T x} \\ e^{\theta_2^T x} \end{bmatrix}$$

利用 softmax regression 回归参数冗余的特点，我们令 $\psi = \theta_1$ ，并且从两个参数向量中都减去向量 θ_1 ，得到：

$$\begin{aligned}
h(x) &= \frac{1}{e^{\vec{0}^T x} + e^{(\theta_2 - \theta_1)^T x}} \begin{bmatrix} e^{\vec{0}^T x} \\ e^{(\theta_1 - \theta_2)^T x} \end{bmatrix} \\
&= \begin{bmatrix} \frac{1}{1 + e^{(\theta_1 - \theta_2)^T x}} \\ \frac{e^{(\theta_1 - \theta_2)^T x}}{1 + e^{(\theta_1 - \theta_2)^T x}} \end{bmatrix} \\
&= \begin{bmatrix} \frac{1}{1 + e^{(\theta_1 - \theta_2)^T x}} \\ 1 - \frac{1}{1 + e^{(\theta_1 - \theta_2)^T x}} \end{bmatrix}
\end{aligned}$$

因此，用 θ' 来表示 $\theta_1 - \theta_2$ ，我们就会发现 softmax regression 预测其中一个类别的概率为 $\frac{1}{1 + e^{(\theta')^T x}}$ ，另一个类别的概率为 $1 - \frac{1}{1 + e^{(\theta')^T x}}$ ，这与 logistic regression 是一致的。

5. Softmax Regression vs. k 个二元分类器

如果你在开发一个音乐分类的应用，需要对 k 种类型的音乐进行识别，那么是选择使用 softmax regression，还是使用 logistic regression 建立 k 个独立的二元分类器呢？

这一选择取决于你的类别之间是否互斥，例如，如果你有四个类别的音乐，分别为：古典音乐、乡村音乐、摇滚乐和爵士乐，那么你可以假设每个训练样本只会被打上一个标签（即一首歌只能属于这四种音乐类型的其中一种），此时，你应该使用类别数 $k = 4$ 的 softmax regression（如果在你的数据集中，有的歌曲不属于以上四类的其中任何一类，那么你可以添加一个“其他类”，并将类别数 k 设为5）。

如果你的四个类别如下：人声音乐、舞曲、影视原声、流行歌曲，那么这些类别之间并不是互斥的。例如：一首歌曲可以来源于影视原声，同时也包含人声。在这种情况下，使用4个二分类的 logistic regression 更为合适。这样每个新的音乐作品，我们的算法可以分别判断它是否属于各个类别。

现在，我们来看一个计算机视觉领域的例子，你的任务是将图像分到三个不同的类别中。（1）假设这三个类

别分别是：室内场景、户外城区场景、户外荒野场景。你会使用softmax regression 还是3个 logistic regression 呢？（2）现在假设这三个类别分别是室内场景、黑白图片、包含人物的图片，你会选择 softmax regression 还是多个 logistic regression 呢？

在第一个例子中，三个类别是互斥的，因此，更适于选择 softmax regression。而在第二个例子中，建立三个独立的 logistic regression 更加合适。

In []: